

openclibration

with Julian Kent

Zurich C++ Meetup

Agenda

- What is photogrammetry
- Typical photogrammetry pipeline
- What makes `openCalibration` different
- What's next for `openCalibration`
- Demo

Photogrammetry

- Images \rightarrow Camera model + sparse 3D point cloud
- Camera model: lens distortions + global position / orientation

Typical Photogrammetry Pipeline

Bundler


- Image loading
- Feature Extraction
- Matching
- RANSAC
- Decomposition, or PnP
- Bundle Adjustment



Typical Photogrammetry Pipeline

Bundler

- Image loading
- Feature Extraction
- Matching
- RANSAC
- PnP
- Bundle Adjustment



The opencalibration pipeline

- Load and initial alignment
- Camera pose & lens parameter refinement
- Low-res surface model estimation
- DSM + ortho preview (in progress)
 - Thumbnails from input images are used to quickly generate a low-resolution orthomosaic
 - Also generate a DSM to go with it

The opencalibration pipeline

Load and initial alignment

- Works 'online', new data can be added while processing

Split into different pipelined stages, on batches of images

Stage 1:


- Image Loading & Feature Extraction
 - IO and FPU intensive, uses lots of memory:

Stage 2:

- Matching, RANSAC, Decomposition
 - POPCNT heavy / $O(N^2)$ on features / future GPU offload:

Stage 3:

- Camera orientation optimization with local cluster
 - Lots of small serial operations, single threaded



The opencalibration pipeline

Camera pose & lens parameter refinement


- Mostly a ‘standard, classic’ bundle adjustment
- What’s special:
 - Split images based on Spectral Clustering of measurement graph
 - Alternate between optimizing camera pose and lens parameters
 - Do pose graph minimization on each cluster separately - use a low-resolution surface mesh instead of individual keypoints.
 - Avoid $O(N^3)$ scaling by breaking into clusters and parallelizing
 - Estimate camera parameters with classic keypoints-style bundle adjustment
 - Avoid $O(N^3)$ scaling by only estimating camera parameters on largest cluster

The opencalibration pipeline

Surface model estimation

- Low resolution surface mesh directly optimized in clusters
- Meshes merged and final global optimization at the end - $O(N^3)$

opencalibration - 2m45s on a laptop



Pix4D - 22min just for calibration



The opencalibration pipeline

DSM + ortho preview (in progress)

- Thumbnails from input images are used to quickly generate a low-resolution orthomosaic
- Also generate a DSM to go with it

What's next for openCalibration

- Quality metrics - need ground truth datasets
- Shifting clusters to avoid discontinuities
 - Weight edges which were 'cut' in previous iterations higher
- Inverse camera model experiments
 - Right now the camera models are defined 3D \rightarrow 2D
 - Makes directly optimizing camera parameters on mesh slow/complicated
 - Idea: define camera model the other way, fit, minimize, fit back
- Making actually useful outputs:
 - Surface models: 3D mesh / Digital Surface Map
 - Orthomosaic
 - Idea: a scaled up version of the thumbnail preview
- Sharing it with the team at OpenDroneMap
 - Still a long way to go before getting here

Questions

<https://github.com/jkflying/opencalibration/>